

June 24, 2024

1 Six Degrees of Kevin Bacon

Introduction - Six Degrees of Kevin Bacon is a game based on the “six degrees of separation” concept, which posits that any two people on Earth are six or fewer acquaintance links apart. Movie buffs challenge each other to find the shortest path between an arbitrary actor and prolific actor Kevin Bacon. It rests on the assumption that anyone involved in the film industry can be linked through their film roles to Bacon within six steps. The analysis of social networks can be a computationally intensive task, especially when dealing with large volumes of data. It is also a challenging problem to devise a correct methodology to infer an informative social network structure. Here, we will analyze a social network of actors and actresses that co-participated in movies. We will do some simple descriptive analysis, and in the end try to relate an actor/actress’s position in the social network with the success of the movies in which they participate.

Rules & Notes - Please take your time to read the following points:

1. The submission deadline shall be set for the 10th of June at 23:59.
2. It is acceptable that you **discuss** with your colleagues different approaches to solve each step of the problem set. You are responsible for writing your own code, and analysing the results. Clear cases of cheating will be penalized with 0 points in this assignment;
3. After review of your submission files, and before a mark is attributed, you might be called to orally defend your submission;
4. You will be scored first and foremost by the number of correct answers, secondly by the logic used in the trying to approach each step of the problem set;
5. Consider skipping questions that you are stuck in, and get back to them later;
6. Expect computations to take a few minutes to finish in some of the steps.
7. **IMPORTANT** It is expected you have developed skills beyond writting SQL queries. Any question where you directly write a SQL query (then for example create a temporary table and use spark.sql to pass the query) will receive a 25% penalty. Using the Spark syntax (for example `dataframe.select("*").where("conditions")`) is acceptable and does not incur this penalty. Comment your code in a reasonable fashion.
8. **Questions** – Any questions about this assignment should be posted in the Forum@Moodle. The last class will be an open office session for anyone with questions concerning the assignment.
9. **Delivery** - To fulfil this activity you will have to upload the following materials to Moodle:
 1. An exported IPython notebook. The notebook should be solved (have results displayed), but should contain all neccesary code so that when the notebook is run in databricks it should also replicate these results. This means the all data downloading and processing should be done in this notebook. It is also important you clearly indicate where your final

answer to each question is when you are using multiple cells (for example you print “my final answer is” before your answer or use cell comments). Please make sure to name your file in the following way: *[student_number1]_[student_number2]_submission.ipynb. *As an example:* 19740001_197400010_submission.ipynb*

2. **Delivery** - You will also need to provide a signed statement of authorship, which is present in the last page;
 3. It is recommended you read the whole assignment before starting.
 4. You can add as many cells as you like to answer the questions.
 5. You can make use of caching or persisting your RDDs or Dataframes, this may speed up performance.
 6. If you have trouble with graphframes in databricks (specifically the import statement) you need to make sure the graphframes package is installed on the cluster you are running. If you click home on the left, then click on the graphframes library, from where you can install the package on your cluster (check the graphframes checkbox and click install). Another installation option is using the JAR available on Moodle with the graphframes library.
10. **Note:** By including the name and student number of each group member in the submission notebook, this will be considered as a declaration of authorship.

Data Sources and Description We will use data from IMDB. You can download raw datafiles from <https://datasets.imdbws.com>. Note that the files are tab delimited (.tsv) You can find a description of the each datafile in <https://www.imdb.com/interfaces/>

1.1 Questions

1.1.1 Data loading and preperation

Review the file descriptions and load the necessary data onto your databricks cluser and into spark dataframes. You will need to use shell commands to download the data, unzip the data, load the data into spark. Note that the data might require parsing and preprocessing to be ready for the questions below.

Hints You can use ‘gunzip’ to unzip the .tz files. The data files will then be tab seperated (.tsv), which you can load into a dataframe using the tab seperated option instead of the comma seperated option we have typically used in class: `.option("sep","\t")`

```
[ ]: %sh
wget "https://datasets.imdbws.com/name.basics.tsv.gz" -O /tmp/name.basics.tsv.gz
wget "https://datasets.imdbws.com/title.akas.tsv.gz" -O /tmp/title.akas.tsv.gz
wget "https://datasets.imdbws.com/title.basics.tsv.gz" -O /tmp/title.basics.tsv.
    ↪gz
wget "https://datasets.imdbws.com/title.crew.tsv.gz" -O /tmp/title.crew.tsv.gz
wget "https://datasets.imdbws.com/title.episode.tsv.gz" -O /tmp/title.episode.
    ↪tsv.gz
wget "https://datasets.imdbws.com/title.principals.tsv.gz" -O /tmp/title.
    ↪principals.tsv.gz
```

```
wget "https://datasets.imdbws.com/title.ratings.tsv.gz" -O /tmp/title.ratings.  
↳tsv.gz
```

```
--2024-06-09 07:53:56-- https://datasets.imdbws.com/name.basics.tsv.gz  
Resolving datasets.imdbws.com (datasets.imdbws.com)... 18.245.253.55,  
18.245.253.117, 18.245.253.85, ...  
Connecting to datasets.imdbws.com (datasets.imdbws.com)|18.245.253.55|:443...  
connected.  
HTTP request sent, awaiting response... 200 OK  
Length: 267637948 (255M) [binary/octet-stream]  
Saving to: '/tmp/name.basics.tsv.gz'
```

```
OK ... .. 0% 15.2M 17s  
50K ... .. 0% 20.8M 15s  
100K ... .. 0% 19.6M 14s  
150K ... .. 0% 21.1M 14s  
200K ... .. 0% 21.8M 13s  
250K ... .. 0% 91.1M 11s  
300K ... .. 0% 68.3M 10s  
350K ... .. 0% 32.4M 10s  
400K ... .. 0% 45.4M 10s  
450K ... .. 0% 30.1M 9s  
500K ... .. 0% 49.4M 9s  
550K ... .. 0% 115M 8s  
600K ... .. 0% 66.6M 8s  
650K ... .. 0% 205M 8s  
700K ... .. 0% 76.7M 7s  
750K ... .. 0% 74.5M 7s  
800K ... .. 0% 114M 7s  
850K ... .. 0% 87.5M 7s  
900K ... .. 0% 76.2M 6s  
950K ... .. 0% 338M 6s  
1000K ... .. 0% 67.4M 6s  
1050K ... .. 0% 87.7M 6s  
1100K ... .. 0% 84.4M 6s  
1150K ... .. 0% 66.8M 6s  
1200K ... .. 0% 368M 5s  
1250K ... .. 0% 91.2M 5s  
1300K ... .. 0% 111M 5s  
1350K ... .. 0% 86.4M 5s  
1400K ... .. 0% 126M 5s  
1450K ... .. 0% 445M 5s  
1500K ... .. 0% 110M 5s  
1550K ... .. 0% 73.7M 5s  
1600K ... .. 0% 138M 5s  
1650K ... .. 0% 109M 5s  
1700K ... .. 0% 87.7M 5s  
1750K ... .. 0% 382M 4s
```

1800K	0%	114M	4s
1850K	0%	120M	4s
1900K	0%	127M	4s
1950K	0%	296M	4s
2000K	0%	103M	4s
2050K	0%	129M	4s
2100K	0%	117M	4s
2150K	0%	105M	4s
2200K	0%	329M	4s
2250K	0%	104M	4s
2300K	0%	181M	4s
2350K	0%	153M	4s
2400K	0%	126M	4s
2450K	0%	213M	4s
2500K	0%	190M	4s
2550K	0%	142M	4s
2600K	1%	117M	4s
2650K	1%	150M	4s
2700K	1%	410M	4s
2750K	1%	111M	3s
2800K	1%	142M	3s
2850K	1%	165M	3s
2900K	1%	151M	3s
2950K	1%	331M	3s
3000K	1%	150M	3s
3050K	1%	156M	3s
3100K	1%	155M	3s
3150K	1%	404M	3s
3200K	1%	137M	3s
3250K	1%	185M	3s
3300K	1%	162M	3s
3350K	1%	168M	3s
3400K	1%	152M	3s
3450K	1%	365M	3s
3500K	1%	172M	3s
3550K	1%	132M	3s
3600K	1%	207M	3s
3650K	1%	538M	3s
3700K	1%	144M	3s
3750K	1%	230M	3s
3800K	1%	150M	3s
3850K	1%	207M	3s
3900K	1%	431M	3s
3950K	1%	154M	3s
4000K	1%	275M	3s
4050K	1%	228M	3s
4100K	1%	203M	3s
4150K	1%	221M	3s

4200K	1%	329M	3s
4250K	1%	196M	3s
4300K	1%	210M	3s
4350K	1%	190M	3s
4400K	1%	219M	3s
4450K	1%	300M	3s
4500K	1%	190M	3s
4550K	1%	264M	3s
4600K	1%	225M	3s
4650K	1%	343M	3s
4700K	1%	198M	3s
4750K	1%	224M	3s
4800K	1%	64.2M	3s
4850K	1%	304M	3s
4900K	1%	363M	3s
4950K	1%	107M	3s
5000K	1%	176M	3s
5050K	1%	388M	2s
5100K	1%	156M	2s
5150K	1%	261M	2s
5200K	2%	564M	2s
5250K	2%	184M	2s
5300K	2%	243M	2s
5350K	2%	158M	2s
5400K	2%	323M	2s
5450K	2%	233M	2s
5500K	2%	573M	2s
5550K	2%	241M	2s
5600K	2%	263M	2s
5650K	2%	390M	2s
5700K	2%	429M	2s
5750K	2%	230M	2s
5800K	2%	250M	2s
5850K	2%	287M	2s
5900K	2%	326M	2s
5950K	2%	272M	2s
6000K	2%	277M	2s
6050K	2%	308M	2s
6100K	2%	308M	2s
6150K	2%	330M	2s
6200K	2%	400M	2s
6250K	2%	298M	2s
6300K	2%	364M	2s
6350K	2%	318M	2s
6400K	2%	391M	2s
6450K	2%	358M	2s
6500K	2%	421M	2s
6550K	2%	347M	2s

6600K	2%	420M	2s
6650K	2%	314M	2s
6700K	2%	427M	2s
6750K	2%	344M	2s
6800K	2%	358M	2s
6850K	2%	410M	2s
6900K	2%	439M	2s
6950K	2%	381M	2s
7000K	2%	358M	2s
7050K	2%	347M	2s
7100K	2%	418M	2s
7150K	2%	390M	2s
7200K	2%	319M	2s
7250K	2%	419M	2s
7300K	2%	396M	2s
7350K	2%	326M	2s
7400K	2%	362M	2s
7450K	2%	348M	2s
7500K	2%	392M	2s
7550K	2%	359M	2s
7600K	2%	308M	2s
7650K	2%	389M	2s
7700K	2%	351M	2s
7750K	2%	364M	2s
7800K	3%	391M	2s
7850K	3%	310M	2s
7900K	3%	394M	2s
7950K	3%	318M	2s
8000K	3%	350M	2s
8050K	3%	398M	2s
8100K	3%	278M	2s
8150K	3%	384M	2s
8200K	3%	343M	2s
8250K	3%	339M	2s
8300K	3%	339M	2s
8350K	3%	359M	2s
8400K	3%	377M	2s
8450K	3%	342M	2s
8500K	3%	368M	2s
8550K	3%	299M	2s
8600K	3%	362M	2s
8650K	3%	273M	2s
8700K	3%	337M	2s
8750K	3%	365M	2s
8800K	3%	355M	2s
8850K	3%	374M	2s
8900K	3%	324M	2s
8950K	3%	362M	2s

9000K	3%	373M	2s
9050K	3%	387M	2s
9100K	3%	316M	2s
9150K	3%	346M	2s
9200K	3%	326M	2s
9250K	3%	431M	2s
9300K	3%	314M	2s
9350K	3%	421M	2s
9400K	3%	387M	2s
9450K	3%	295M	2s
9500K	3%	341M	2s
9550K	3%	345M	2s
9600K	3%	314M	2s
9650K	3%	335M	2s
9700K	3%	388M	2s
9750K	3%	401M	2s
9800K	3%	366M	2s
9850K	3%	366M	2s
9900K	3%	393M	2s
9950K	3%	383M	2s
10000K	3%	284M	2s
10050K	3%	388M	2s
10100K	3%	296M	2s
10150K	3%	343M	2s
10200K	3%	368M	2s
10250K	3%	299M	2s
10300K	3%	344M	2s
10350K	3%	323M	2s
10400K	3%	354M	2s
10450K	4%	308M	2s
10500K	4%	331M	2s
10550K	4%	336M	2s
10600K	4%	329M	2s
10650K	4%	325M	2s
10700K	4%	345M	2s
10750K	4%	259M	2s
10800K	4%	319M	2s
10850K	4%	341M	2s
10900K	4%	353M	2s
10950K	4%	311M	2s
11000K	4%	265M	2s
11050K	4%	337M	2s
11100K	4%	327M	2s
11150K	4%	335M	2s
11200K	4%	277M	2s
11250K	4%	330M	1s
11300K	4%	337M	1s
11350K	4%	332M	1s

11400K	4%	283M	1s
11450K	4%	295M	1s
11500K	4%	323M	1s
11550K	4%	353M	1s
11600K	4%	344M	1s
11650K	4%	282M	1s
11700K	4%	341M	1s
11750K	4%	349M	1s
11800K	4%	380M	1s
11850K	4%	302M	1s
11900K	4%	315M	1s
11950K	4%	383M	1s
12000K	4%	339M	1s
12050K	4%	305M	1s
12100K	4%	360M	1s
12150K	4%	337M	1s
12200K	4%	410M	1s
12250K	4%	299M	1s
12300K	4%	347M	1s
12350K	4%	327M	1s
12400K	4%	370M	1s
12450K	4%	322M	1s
12500K	4%	326M	1s
12550K	4%	304M	1s
12600K	4%	329M	1s
12650K	4%	319M	1s
12700K	4%	310M	1s
12750K	4%	247M	1s
12800K	4%	350M	1s
12850K	4%	360M	1s
12900K	4%	399M	1s
12950K	4%	487M	1s
13000K	4%	330M	1s
13050K	5%	365M	1s
13100K	5%	347M	1s
13150K	5%	327M	1s
13200K	5%	317M	1s
13250K	5%	273M	1s
13300K	5%	409M	1s
13350K	5%	372M	1s
13400K	5%	326M	1s
13450K	5%	412M	1s
13500K	5%	389M	1s
13550K	5%	377M	1s
13600K	5%	367M	1s
13650K	5%	314M	1s
13700K	5%	394M	1s
13750K	5%	425M	1s

13800K	5%	327M	1s
13850K	5%	356M	1s
13900K	5%	308M	1s
13950K	5%	362M	1s
14000K	5%	289M	1s
14050K	5%	417M	1s
14100K	5%	440M	1s
14150K	5%	321M	1s
14200K	5%	280M	1s
14250K	5%	315M	1s
14300K	5%	362M	1s
14350K	5%	326M	1s
14400K	5%	326M	1s
14450K	5%	315M	1s
14500K	5%	358M	1s
14550K	5%	316M	1s
14600K	5%	359M	1s
14650K	5%	277M	1s
14700K	5%	330M	1s
14750K	5%	275M	1s
14800K	5%	351M	1s
14850K	5%	346M	1s
14900K	5%	272M	1s
14950K	5%	298M	1s
15000K	5%	337M	1s
15050K	5%	310M	1s
15100K	5%	347M	1s
15150K	5%	279M	1s
15200K	5%	336M	1s
15250K	5%	357M	1s
15300K	5%	305M	1s
15350K	5%	331M	1s
15400K	5%	321M	1s
15450K	5%	338M	1s
15500K	5%	319M	1s
15550K	5%	320M	1s
15600K	5%	384M	1s
15650K	6%	322M	1s
15700K	6%	411M	1s
15750K	6%	424M	1s
15800K	6%	339M	1s
15850K	6%	365M	1s
15900K	6%	334M	1s
15950K	6%	361M	1s
16000K	6%	342M	1s
16050K	6%	278M	1s
16100K	6%	311M	1s
16150K			

*** WARNING: max output size exceeded, skipping output. ***

... 98% 341M 0s

643750K	98%	1.92M	0s
643800K	98%	317M	0s
643850K	98%	390M	0s
643900K	98%	380M	0s
643950K	98%	328M	0s
644000K	98%	386M	0s
644050K	98%	367M	0s
644100K	98%	273M	0s
644150K	98%	353M	0s
644200K	98%	386M	0s
644250K	98%	339M	0s
644300K	98%	371M	0s
644350K	98%	377M	0s
644400K	98%	319M	0s
644450K	98%	378M	0s
644500K	98%	392M	0s
644550K	98%	368M	0s
644600K	98%	318M	0s
644650K	98%	351M	0s
644700K	98%	339M	0s
644750K	98%	343M	0s
644800K	98%	370M	0s
644850K	98%	367M	0s
644900K	98%	323M	0s
644950K	98%	431M	0s
645000K	98%	393M	0s
645050K	98%	388M	0s
645100K	98%	345M	0s
645150K	98%	377M	0s
645200K	98%	367M	0s
645250K	98%	454M	0s
645300K	98%	385M	0s
645350K	98%	351M	0s
645400K	98%	342M	0s
645450K	98%	377M	0s
645500K	98%	345M	0s
645550K	98%	365M	0s
645600K	98%	256M	0s
645650K	98%	384M	0s
645700K	98%	315M	0s
645750K	98%	339M	0s
645800K	98%	378M	0s
645850K	98%	343M	0s
645900K	98%	383M	0s

645950K	98%	378M	0s
646000K	98%	355M	0s
646050K	98%	372M	0s
646100K	99%	359M	0s
646150K	99%	397M	0s
646200K	99%	359M	0s
646250K	99%	325M	0s
646300K	99%	368M	0s
646350K	99%	357M	0s
646400K	99%	346M	0s
646450K	99%	348M	0s
646500K	99%	384M	0s
646550K	99%	359M	0s
646600K	99%	380M	0s
646650K	99%	323M	0s
646700K	99%	387M	0s
646750K	99%	378M	0s
646800K	99%	315M	0s
646850K	99%	375M	0s
646900K	99%	352M	0s
646950K	99%	374M	0s
647000K	99%	383M	0s
647050K	99%	309M	0s
647100K	99%	368M	0s
647150K	99%	331M	0s
647200K	99%	363M	0s
647250K	99%	367M	0s
647300K	99%	389M	0s
647350K	99%	382M	0s
647400K	99%	393M	0s
647450K	99%	408M	0s
647500K	99%	455M	0s
647550K	99%	338M	0s
647600K	99%	424M	0s
647650K	99%	354M	0s
647700K	99%	375M	0s
647750K	99%	340M	0s
647800K	99%	347M	0s
647850K	99%	310M	0s
647900K	99%	409M	0s
647950K	99%	381M	0s
648000K	99%	332M	0s
648050K	99%	267M	0s
648100K	99%	360M	0s
648150K	99%	381M	0s
648200K	99%	367M	0s
648250K	99%	276M	0s
648300K	99%	371M	0s

648350K	99%	344M	0s
648400K	99%	389M	0s
648450K	99%	305M	0s
648500K	99%	285M	0s
648550K	99%	324M	0s
648600K	99%	363M	0s
648650K	99%	300M	0s
648700K	99%	369M	0s
648750K	99%	335M	0s
648800K	99%	309M	0s
648850K	99%	367M	0s
648900K	99%	282M	0s
648950K	99%	370M	0s
649000K	99%	378M	0s
649050K	99%	376M	0s
649100K	99%	1.17M	0s
649150K	99%	466M	0s
649200K	99%	356M	0s
649250K	99%	511M	0s
649300K	99%	466M	0s
649350K	99%	419M	0s
649400K	99%	381M	0s
649450K	99%	433M	0s
649500K	99%	346M	0s
649550K	99%	445M	0s
649600K	99%	444M	0s
649650K	99%	486M	0s
649700K	99%	425M	0s
649750K	99%	474M	0s
649800K	99%	476M	0s
649850K	99%	497M	0s
649900K	99%	421M	0s
649950K	99%	474M	0s
650000K	99%	392M	0s
650050K	99%	467M	0s
650100K	99%	377M	0s
650150K	99%	442M	0s
650200K	99%	462M	0s
650250K	99%	444M	0s
650300K	99%	451M	0s
650350K	99%	393M	0s
650400K	99%	471M	0s
650450K	99%	390M	0s
650500K	99%	459M	0s
650550K	99%	477M	0s
650600K	99%	367M	0s
650650K	99%	452M	0s
650700K	99%	444M	0s

650750K	...	99%	445M	0s
650800K	...	99%	474M	0s
650850K	...	99%	386M	0s
650900K	...	99%	566M	0s
650950K	...	99%	518M	0s
651000K	...	99%	436M	0s
651050K	...	99%	440M	0s
651100K	...	99%	490M	0s
651150K	...	99%	433M	0s
651200K	...	99%	411M	0s
651250K	...	99%	322M	0s
651300K	...	99%	267M	0s
651350K	...	99%	258M	0s
651400K	...	99%	311M	0s
651450K	...	99%	381M	0s
651500K	...	99%	611M	0s
651550K	...	99%	623M	0s
651600K	...	99%	549M	0s
651650K	...	99%	627M	0s
651700K	...	99%	557M	0s
651750K	...	99%	606M	0s
651800K	...	99%	650M	0s
651850K	...	99%	560M	0s
651900K	...	99%	635M	0s
651950K	...	99%	574M	0s
652000K	...	99%	617M	0s
652050K	...	99%	537M	0s
652100K	...	99%	569M	0s
652150K	...	99%	635M	0s
652200K	...	99%	569M	0s
652250K	...	99%	631M	0s
652300K	...	99%	586M	0s
652350K	...	99%	559M	0s
652400K	...	99%	637M	0s
652450K	...	99%	585M	0s
652500K	...	99%	570M	0s
652550K	...	99%	575M	0s
652600K	...	99%	445M	0s
652650K	...			

100% 558M=6.5s

2024-06-09 07:54:12 (98.6 MB/s) - '/tmp/title.principals.tsv.gz' saved
[668327000/668327000]

--2024-06-09 07:54:12-- <https://datasets.imdbws.com/title.ratings.tsv.gz>
Resolving datasets.imdbws.com (datasets.imdbws.com)... 18.245.253.55,
18.245.253.70, 18.245.253.85, ...
Connecting to datasets.imdbws.com (datasets.imdbws.com)|18.245.253.55|:443...
connected.

HTTP request sent, awaiting response... 200 OK
Length: 7261619 (6.9M) [binary/octet-stream]
Saving to: '/tmp/title.ratings.tsv.gz'

0K	0%	21.0M	0s
50K	1%	17.8M	0s
100K	2%	22.0M	0s
150K	2%	108M	0s
200K	3%	59.1M	0s
250K	4%	87.5M	0s
300K	4%	59.4M	0s
350K	5%	166M	0s
400K	6%	83.3M	0s
450K	7%	68.3M	0s
500K	7%	180M	0s
550K	8%	109M	0s
600K	9%	327M	0s
650K	9%	73.5M	0s
700K	10%	301M	0s
750K	11%	248M	0s
800K	11%	329M	0s
850K	12%	69.0M	0s
900K	13%	77.8M	0s
950K	14%	85.6M	0s
1000K	14%	276M	0s
1050K	15%	243M	0s
1100K	16%	60.0M	0s
1150K	16%	244M	0s
1200K	17%	153M	0s
1250K	18%	364M	0s
1300K	19%	329M	0s
1350K	19%	296M	0s
1400K	20%	371M	0s
1450K	21%	4.32M	0s
1500K	21%	251M	0s
1550K	22%	270M	0s
1600K	23%	300M	0s
1650K	23%	271M	0s
1700K	24%	301M	0s
1750K	25%	246M	0s
1800K	26%	295M	0s
1850K	26%	302M	0s
1900K	27%	311M	0s
1950K	28%	282M	0s
2000K	28%	273M	0s
2050K	29%	313M	0s
2100K	30%	335M	0s
2150K	31%	251M	0s

2200K	31%	87.8M	0s
2250K	32%	102M	0s
2300K	33%	295M	0s
2350K	33%	333M	0s
2400K	34%	310M	0s
2450K	35%	169M	0s
2500K	35%	123M	0s
2550K	36%	74.7M	0s
2600K	37%	124M	0s
2650K	38%	108M	0s
2700K	38%	291M	0s
2750K	39%	23.3M	0s
2800K	40%	320M	0s
2850K	40%	193M	0s
2900K	41%	323M	0s
2950K	42%	332M	0s
3000K	43%	285M	0s
3050K	43%	327M	0s
3100K	44%	290M	0s
3150K	45%	340M	0s
3200K	45%	294M	0s
3250K	46%	353M	0s
3300K	47%	334M	0s
3350K	47%	322M	0s
3400K	48%	335M	0s
3450K	49%	257M	0s
3500K	50%	291M	0s
3550K	50%	315M	0s
3600K	51%	332M	0s
3650K	52%	345M	0s
3700K	52%	368M	0s
3750K	53%	368M	0s
3800K	54%	328M	0s
3850K	54%	369M	0s
3900K	55%	401M	0s
3950K	56%	286M	0s
4000K	57%	321M	0s
4050K	57%	311M	0s
4100K	58%	324M	0s
4150K	59%	284M	0s
4200K	59%	334M	0s
4250K	60%	311M	0s
4300K	61%	322M	0s
4350K	62%	306M	0s
4400K	62%	302M	0s
4450K	63%	302M	0s
4500K	64%	331M	0s
4550K	64%	329M	0s

4600K	65%	318M	0s
4650K	66%	325M	0s
4700K	66%	328M	0s
4750K	67%	276M	0s
4800K	68%	296M	0s
4850K	69%	292M	0s
4900K	69%	301M	0s
4950K	70%	301M	0s
5000K	71%	304M	0s
5050K	71%	308M	0s
5100K	72%	287M	0s
5150K	73%	125M	0s
5200K	74%	179M	0s
5250K	74%	287M	0s
5300K	75%	292M	0s
5350K	76%	291M	0s
5400K	76%	322M	0s
5450K	77%	296M	0s
5500K	78%	325M	0s
5550K	78%	266M	0s
5600K	79%	61.9M	0s
5650K	80%	237M	0s
5700K	81%	194M	0s
5750K	81%	254M	0s
5800K	82%	289M	0s
5850K	83%	213M	0s
5900K	83%	225M	0s
5950K	84%	293M	0s
6000K	85%	335M	0s
6050K	86%	292M	0s
6100K	86%	281M	0s
6150K	87%	334M	0s
6200K	88%	392M	0s
6250K	88%	325M	0s
6300K	89%	502M	0s
6350K	90%	317M	0s
6400K	90%	402M	0s
6450K	91%	328M	0s
6500K	92%	273M	0s
6550K	93%	350M	0s
6600K	93%	343M	0s
6650K	94%	283M	0s
6700K	95%	308M	0s
6750K	95%	307M	0s
6800K	96%	333M	0s
6850K	97%	1.12M	0s
6900K	98%	238M	0s
6950K	98%	445M	0s


```
7000K ... .. 99% 358M 0s
7050K ... .. . 100% 267M=0.1s
```

```
2024-06-09 07:54:12 (72.5 MB/s) - '/tmp/title.ratings.tsv.gz' saved
[7261619/7261619]
```

```
[ ]: %sh
```

```
gunzip -f /tmp/name.basics.tsv.gz
gunzip -f /tmp/title.akas.tsv.gz
gunzip -f /tmp/title.basics.tsv.gz
gunzip -f /tmp/title.crew.tsv.gz
gunzip -f /tmp/title.episode.tsv.gz
gunzip -f /tmp/title.principals.tsv.gz
gunzip -f /tmp/title.ratings.tsv.gz
```

```
[ ]: dbutils.fs.mkdirs("/mnt/data/test")
```

```
Out[3]: True
```

```
[ ]: dbutils.fs.mv("file:/tmp/name.basics.tsv", "dbfs:/mnt/data/test")
dbutils.fs.mv("file:/tmp/title.akas.tsv", "dbfs:/mnt/data/test")
dbutils.fs.mv("file:/tmp/title.basics.tsv", "dbfs:/mnt/data/test")
dbutils.fs.mv("file:/tmp/title.crew.tsv", "dbfs:/mnt/data/test")
dbutils.fs.mv("file:/tmp/title.episode.tsv", "dbfs:/mnt/data/test")
dbutils.fs.mv("file:/tmp/title.principals.tsv", "dbfs:/mnt/data/test")
dbutils.fs.mv("file:/tmp/title.ratings.tsv", "dbfs:/mnt/data/test")
```

```
Out[4]: True
```

```
[ ]: name_basics_df = spark.read.option('inferSchema', 'true').option('header', ␣
    ↪ 'true').option('sep', '\t').csv('dbfs:/mnt/data/test/name.basics.tsv')
title_akas_df = spark.read.option('inferSchema', 'true').option('header', ␣
    ↪ 'true').option('sep', '\t').csv('dbfs:/mnt/data/test/title.akas.tsv')
title_basics_df = spark.read.option('inferSchema', 'true').option('header', ␣
    ↪ 'true').option('sep', '\t').csv('dbfs:/mnt/data/test/title.basics.tsv')
title_crew_df = spark.read.option('inferSchema', 'true').option('header', ␣
    ↪ 'true').option('sep', '\t').csv('dbfs:/mnt/data/test/title.crew.tsv')
title_episode_df = spark.read.option('inferSchema', 'true').option('header', ␣
    ↪ 'true').option('sep', '\t').csv('dbfs:/mnt/data/test/title.episode.tsv')
title_principals_df = spark.read.option('inferSchema', 'true').option('header', ␣
    ↪ 'true').option('sep', '\t').csv('dbfs:/mnt/data/test/title.principals.tsv')
title_ratings_df = spark.read.option('inferSchema', 'true').option('header', ␣
    ↪ 'true').option('sep', '\t').csv('dbfs:/mnt/data/test/title.ratings.tsv')
```

1.1.2 Network Inference, Let's build a network

In the following questions you will look to summarise the data and build a network. We want to examine a network that abstracts how actors and actress are related through their co-participation in movies. To that end perform the following steps:

Q1 Create a DataFrame that combines **all the information** on each of the titles (i.e., movies, tv-shows, etc ...) and **all of the information** the participants in those movies (i.e., actors, directors, etc ...), make sure the actual names of the movies and participants are included. It may be worth reviewing the following questions to see how this dataframe will be used.

How many rows does your dataframe have?

```
[ ]: from pyspark.sql.functions import col

[ ]: # Aliasing DataFrames for clarity
titles_df = title_basics_df.alias("titles")
principals_df = title_principals_df.alias("principals")
names_df = name_basics_df.alias("names")
ratings_df = title_ratings_df.alias("ratings")

# Join titles with principals on tconst
titles_with_principals = titles_df.join(
    principals_df,
    col("titles.tconst") == col("principals.tconst"),
    "left"
).select(
    col("titles.tconst").alias("title_tconst"),
    col("titles.*"),
    col("principals.*")
)

# Join the result with names on nconst
titles_principals_names = titles_with_principals.join(
    names_df,
    col("principals.nconst") == col("names.nconst"),
    "left_outer"
).select(
    col("title_tconst"),
    col("titles.*"),
    col("principals.nconst").alias("principal_nconst"),
    col("principals.*"),
    col("names.*")
)

# Join with ratings on tconst
full_title_info = titles_principals_names.join(
    ratings_df,
    col("title_tconst") == col("ratings.tconst"),
```

```

    "left"
).select(
    col("title_tconst"),
    col("titles.*"),
    col("principal_nconst"),
    col("names.nconst").alias("name_nconst"),
    col("names.*"),
    col("ratings.*"),
    col("principals.*")
)

```

Final selection to avoid duplicate columns

```

final_df = full_title_info.select(
    col("title_tconst").alias("tconst"),
    col("principal_nconst"),
    col("name_nconst"),
    col("ratings.averageRating"),
    col("ratings.numVotes"),
    col("names.primaryName"),
    col("names.birthYear"),
    col("names.deathYear"),
    col("names.primaryProfession"),
    col("titles.primaryTitle"),
    col("titles.originalTitle"),
    col("titles.isAdult"),
    col("titles.startYear"),
    col("titles.endYear"),
    col("titles.runtimeMinutes"),
    col("titles.genres"),
    col("principals.category"),
    col("principals.job"),
    col("principals.characters")
)

```

Show some of the data to verify correctness

```

final_df.show(5)

```

```

+-----+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+-----+-----+
+-----+
|  tconst|principal_nconst|name_nconst|averageRating|numVotes|
primaryName|birthYear|deathYear|  primaryProfession|          primaryTitle|
originalTitle|isAdult|startYear|endYear|runtimeMinutes|          genres|category|
job|          characters|
+-----+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+-----+-----+

```

```

+-----+
|tt0000658|      nm0169871| nm0169871|      6.4|      292|      Émile
Cohl|      1857|      1938|director,animatio...|The Puppet's Nigh...|Le cauchemar
de F...|      0|      1908|      \N|      2|Animation,Short|director|
\N|      \N|
|tt0000839|      nm0294276| nm0294276|      null|      null|      Theo
Frenkel|      1871|      1956|director,actor,wr...| The Curse of Money| The
Curse of Money|      0|      1909|      \N|      \N|
Drama,Short|director|director|      \N|
|tt0000839|      nm0378408| nm0378408|      null|      null| Cecil M.
Hepworth|      1873|      1953|producer,cinemat...| The Curse of Money| The
Curse of Money|      0|      1909|      \N|      \N|
Drama,Short|producer|producer|      \N|
|tt0001170|      nm1400009| nm1400009|      null|      null|William A.
Russell|      1878|      1914|      actor|A Cowboy's Vindic...|A Cowboy's
Vindic...|      0|      1910|      \N|      \N| Short,Western| actor|
\N|      \N|
|tt0001170|      nm0355582| nm0355582|      null|      null| Franklyn
Hall|      1886|      \N|actor,writer,dire...|A Cowboy's Vindic...|A Cowboy's
Vindic...|      0|      1910|      \N|      \N| Short,Western| actor|
\N|["Will Morrison"]|
+-----+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+-----+-----+
+-----+
only showing top 5 rows

```

```
[ ]: # Q1 FINAL ANSWER
```

```

# Show the number of rows
print(f"My dataframe has {final_df.count()} rows.")

```

My dataframe has 87282924 rows.

Q2 Create a new DataFrame based on the previous step, with the following removed: 1. Any participant that is not an actor or actress (as measured by the category column); 1. All adult movies; 1. All dead actors or actresses; 1. All actors or actresses born before 1920 or with no date of birth listed; 1. All titles that are not of the type movie.

How many rows does your dataframe have?

```

[ ]: # Filter for actors and actresses
filtered_df = final_df.filter(col("category").isin(["actor", "actress"]))

# Filter out adult movies
filtered_df = filtered_df.filter(col("isAdult") == 0)

# Filter out dead actors/actresses

```

```

filtered_df = filtered_df.filter(col("deathYear")=="\N")

# Filter out actors/actresses born before 1920 or with no date of birth listed
filtered_df = filtered_df.filter((col("birthYear") >= 1920) & (col("birthYear")_
↳ != "\N"))

# Filter for movie titles only
filtered_df = filtered_df.filter(col("titleType") == "movie")

```

```

[ ]: # Q2 FINAL ANSWER

# Show the number of rows
print(f"The filtered dataframe has {filtered_df.count()} rows.")

```

The filtered dataframe has 930698 rows.

Q3 Convert the above Dataframe to an RDD. Use map and reduce to create a paired RDD which counts how many movies each actor / actress appears in.

Display names of the top 10 actors/actresses according to the number of movies in which they appeared. Be careful to deal with different actors / actresses with the same name, these could be different people.

```

[ ]: # Convert DataFrame to RDD
filtered_rdd = filtered_df.rdd

[ ]: # Create a paired RDD (actor_id, 1) and reduce by key to count the movies
actor_movie_counts_rdd = filtered_rdd.map(lambda row: (row['name_nconst'], 1)).
↳ reduceByKey(lambda a, b: a + b)

[ ]: # Convert the filtered DataFrame to RDD to get actor names
actor_names_rdd = filtered_df.select("name_nconst", "primaryName").distinct().
↳ rdd.map(lambda row: (row.name_nconst, row.primaryName))

[ ]: # Join actor_movie_rdd with actor_names_rdd
actor_movie_count_with_names_rdd = actor_movie_counts_rdd.join(actor_names_rdd)

[ ]: # Sort by movie count in descending order and take the top 10
top_10_actors_rdd = actor_movie_count_with_names_rdd.sortBy(lambda x: x[1][0],_
↳ ascending=False).take(10)

[ ]: # Q3 FINAL ANSWER

# Display names of the top 10 actors/actresses
for actor, (count, name) in top_10_actors_rdd:
    print(f"Actor: {name}, Movie Count: {count}")

```

Actor: Brahmanandam, Movie Count: 1130

Actor: Jagathy Sreekumar, Movie Count: 659

Actor: Shakti Kapoor, Movie Count: 600
 Actor: Eric Roberts, Movie Count: 492
 Actor: Aruna Irani, Movie Count: 467
 Actor: Nassar, Movie Count: 440
 Actor: Mammootty, Movie Count: 437
 Actor: Helen, Movie Count: 433
 Actor: Tanikella Bharani, Movie Count: 412
 Actor: Anupam Kher, Movie Count: 409

Q4 Start with the dataframe from Q2. Generate a DataFrame that lists all links of your network. Here we shall consider that a link connects a pair of actors/actresses if they participated in at least one movie together (actors / actresses should be represented by their unique ID's). For every link we then need anytime a pair of actors were together in a movie as a link in each direction (A -> B and B -> A). However links should be distinct we do not need duplicates when two actors worked together in several movies.

Display a DataFrame with the first 10 edges.

```
[ ]: from pyspark.sql.functions import col, collect_list, explode

# Filtered DataFrame from Q2
df_q2 = filtered_df

[ ]: # Select necessary columns
actor_movie_df = filtered_df.select("tconst", "name_nconst")

# Group by movie to get the list of actors in each movie
movie_actors_df = actor_movie_df.groupBy("tconst").
    ↪agg(collect_list("name_nconst").alias("actors"))

# Explode the list of actors to create actor pairs
# First create a new DataFrame where each row contains the movie ID and each
    ↪pair of actors
def create_pairs(actors):
    pairs = []
    for i in range(len(actors)):
        for j in range(i + 1, len(actors)):
            pairs.append((actors[i], actors[j]))
            pairs.append((actors[j], actors[i]))
    return pairs

# Register the function as a UDF
from pyspark.sql.types import ArrayType, StructType, StructField, StringType
from pyspark.sql.functions import udf

pair_schema = ArrayType(StructType([
    StructField("src", StringType(), False),
    StructField("dst", StringType(), False)
```

```

]))

create_pairs_udf = udf(create_pairs, pair_schema)

# Apply the UDF to create pairs
actor_pairs_df = movie_actors_df.withColumn("pairs",
↳explode(create_pairs_udf(col("actors"))))

# Separate src and dst
bidirectional_pairs_df = actor_pairs_df.select(col("pairs.src").alias("src"),
↳col("pairs.dst").alias("dst")).distinct()

```

```

[ ]: # Q4 FINAL ANSWER

# Show the first 10 edges
bidirectional_pairs_df.show(10)

```

```

+-----+-----+
|      src|      dst|
+-----+-----+
|nm0180228|nm0639684|
|nm0172237|nm0149883|
|nm0231942|nm0639684|
|nm0013690|nm0013672|
|nm0180228|nm0231942|
|nm0639684|nm0231942|
|nm0231942|nm0180228|
|nm0571763|nm0177320|
|nm0639684|nm0180228|
|nm0013672|nm0013690|
+-----+-----+
only showing top 10 rows

```

Q5 Compute the page rank of each actor. This can be done using GraphFrames or by using RDDs and the iterative implementation of the PageRank algorithm. Do not take more than 5 iterations and use reset probability = 0.1.

List the top 10 actors / actresses by pagerank.

```

[ ]: from graphframes import GraphFrame
    from pyspark.sql.functions import col

[ ]: # Create vertices DataFrame
vertices_df = filtered_df.select(col("name_nconst").alias("id")).distinct()

[ ]: # Create edges DataFrame from bidirectional_pairs_df
edges_df = bidirectional_pairs_df.select(col("src").alias("src"), col("dst").
↳alias("dst"))

```

```
[ ]: # Initialize GraphFrame
g = GraphFrame(vertices_df, edges_df)
```

```
/databricks/spark/python/pyspark/sql/dataframe.py:170: UserWarning:
DataFrame.sql_ctx is an internal property, and will be removed in future
releases. Use DataFrame.sparkSession instead.
  warnings.warn(
```

```
[ ]: # Compute PageRank
pagerank_results = g.pageRank(resetProbability=0.1, maxIter=5)
```

```
/databricks/spark/python/pyspark/sql/dataframe.py:149: UserWarning: DataFrame
constructor is internal. Do not directly use it.
  warnings.warn("DataFrame constructor is internal. Do not directly use it.")
```

```
[ ]: # Join PageRank results with filtered_df to get actor names
pagerank_with_names = pagerank_results.vertices.join(filtered_df,
↳ pagerank_results.vertices.id == filtered_df.name_nconst, "inner") \
                                                    .select("id", "pagerank",
↳ "primaryName") \
                                                    .distinct()
```

```
[ ]: # Q5 FINAL ANSWER

# Display top 10 actors by PageRank
top_10_actors = pagerank_with_names.orderBy(col("pagerank").desc()).
↳ select("id", "primaryName", "pagerank").limit(10)
top_10_actors.show()
```

```
+-----+-----+-----+
|      id|  primaryName|  pagerank|
+-----+-----+-----+
|nm0000616|    Eric Roberts|62.777102793531824|
|nm0000514| Michael Madsen| 33.85372574721415|
|nm0001803|    Danny Trejo|26.530499578166594|
|nm0202966|    Keith David|24.796193603716564|
|nm0001595|    Michael Paré|24.302493820642567|
|nm0261724|    Joe Estevez|23.867127734262173|
|nm0726223| Richard Riehle| 22.93143628629502|
|nm0000532| Malcolm McDowell| 22.83158457241578|
|nm0442207|    Lloyd Kaufman| 22.66740153494526|
|nm0000448| Lance Henriksen|22.253312809786298|
+-----+-----+-----+
```

Q6: Create an RDD with the number of outDegrees for each actor. Display the top 10 by outdegrees.


```
[ ]: # Convert DataFrame to RDD
pairs_rdd = bidirectional_pairs_df.rdd

[ ]: # Map to paired RDD with (actor_id, 1) for each out-degree connection
outdegrees_rdd = pairs_rdd.map(lambda row: (row.src, 1)).reduceByKey(lambda a, b: a + b)

[ ]: # Convert the filtered DataFrame to RDD to get actor names
actor_names_rdd = filtered_df.select("name_nconst", "primaryName").distinct().
    rdd.map(lambda row: (row.name_nconst, row.primaryName))

[ ]: # Join outDegrees with actor names
outdegrees_with_names_rdd = outdegrees_rdd.join(actor_names_rdd)

[ ]: # Sort by outDegrees in descending order
top_10_outdegrees = outdegrees_with_names_rdd.sortBy(lambda x: x[1][0],
    ascending=False).take(10)

[ ]: # Q6 FINAL ANSWER

# Display the top 10 actors by outDegrees with names
for actor, (outdegree, name) in top_10_outdegrees:
    print(f"Actor: {name}, OutDegrees: {outdegree}")
```

```
Actor: Eric Roberts, OutDegrees: 1338
Actor: Michael Madsen, OutDegrees: 842
Actor: Anupam Kher, OutDegrees: 761
Actor: Keith David, OutDegrees: 708
Actor: Renji Ishibashi, OutDegrees: 704
Actor: Nassar, OutDegrees: 689
Actor: Gérard Depardieu, OutDegrees: 678
Actor: Danny Trejo, OutDegrees: 664
Actor: Akira Emoto, OutDegrees: 659
Actor: Prakash Raj, OutDegrees: 649
```

1.1.3 Let's play Kevin's own game

Q7 Start with the graphframe / dataframe you developed in the previous questions. Using Spark GraphFrame and/or Spark Core library perform the following steps:

1. Identify the id of Kevin Bacon, there are two actors named 'Kevin Bacon', we will use the one with the highest degree, that is, the one that participated in most titles;
2. Estimate the shortest path between every actor in the database actors and Kevin Bacon, keep a dataframe with this information as you will need it later;
3. Summarise the data, that is, count the number of actors at each number of degrees from Kevin Bacon (you will need to deal with actors unconnected to Kevin Bacon, if not connected to Kevin Bacon given these actors / actresses a score/degree of 20).

```
[ ]: from graphframes import GraphFrame
from pyspark.sql.functions import col, explode, when, lit

# Identify the ID of Kevin Bacon with the highest degree
filtered_df.filter(col("primaryName") == "Kevin Bacon").show()
```

```
+-----+-----+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+-----+-----+-----+
+-----+
|      tconst|principal_nconst|name_nconst|averageRating|numVotes|primaryName|birthYear|deathYear|primaryProfession|primaryTitle|originalTitle|isAdult|startYear|endYear|runtimeMinutes|genres|category|job|characters|
+-----+-----+-----+-----+-----+-----+-----+
+-----+-----+-----+-----+-----+-----+-----+
+-----+
| tt0373450|      nm0000102| nm0000102|          6.4|   18805|Kevin Bacon|1958|      \N|actor,producer,di...|Where the Truth Lies|Where the Truth Lies|0|    2005|      \N|    107|Crime,Mystery,Thr...|actor| \N|["Lanny"]|
| tt0119896|      nm0000102| nm0000102|          5.5|   21522|Kevin Bacon|1958|      \N|actor,producer,di...|Picture Perfect|Picture Perfect|0|    1997|      \N|    101|Comedy,Drama,Romance|actor| \N|["Sam"]|
| tt0361127|      nm0000102| nm0000102|          7.1|   35445|Kevin Bacon|1958|      \N|actor,producer,di...|The Woodsman|The Woodsman|0|    2004|      \N|     87|Drama|actor| \N|["Walter"]|
| tt14502344|      nm0000102| nm0000102|          4.0|   11597|Kevin Bacon|1958|      \N|actor,producer,di...|They/Them|They/Them|0|    2022|      \N|    104|Drama,Horror,Mystery|actor| \N|["Owen"]|
| tt0164181|      nm0000102| nm0000102|          6.9|   87830|Kevin Bacon|1958|      \N|actor,producer,di...|Stir of Echoes|Stir of Echoes|0|    1999|      \N|     99|Horror,Mystery,Th...|actor| \N|["Tom"]|
| tt6317762|      nm0000102| nm0000102|          5.5|    1369|Kevin Bacon|1958|      \N|actor,producer,di...|Space Oddity|Space Oddity|0|    2022|      \N|     92|Comedy,Romance,Sc...|actor| \N| ["Jeff McAllister"]|
| tt0093403|      nm0000102| nm0000102|          6.3|     340|Kevin Bacon|1958|      \N|actor,producer,di...|Lemon Sky|Lemon Sky|0|    1988|      \N|    106|Drama|actor| \N|["Alan"]|
| tt0790736|      nm0000102| nm0000102|          5.6|  144006|Kevin Bacon|1958|      \N|actor,producer,di...|R.I.P.D.|R.I.P.D.|
```

0	2013	\N	96	Action,Adventure,...	actor \N
["Hayes"]					
	tt13075730	nm0000102	nm0000102	4.1	1864 Kevin Bacon
1958	\N	actor,producer,di...	One Way	One Way	
0	2022	\N	95	Action,Thriller	actor \N ["Fred
Sullivan S...					
	tt1512235	nm0000102	nm0000102	6.7	83837 Kevin Bacon
1958	\N	actor,producer,di...	Super	Super	
0	2010	\N	96	Action,Comedy,Crime	actor \N
["Jacques"]					
	tt1578882	nm0000102	nm0000102	5.0	11115 Kevin Bacon
1958	\N	actor,producer,di...	Elephant White	Elephant White	
0	2011	\N	91	Action,Crime,Thri...	actor \N
["Jimmy"]					
	tt8201852	nm0000102	nm0000102	5.4	26106 Kevin Bacon
1958	\N	actor,producer,di...	You Should Have Left	You Should Have Left	
0	2020	\N	93	Horror,Mystery,Th...	actor \N
["Theo"]					
	tt8201852	nm0000102	nm0000102	5.4	26106 Kevin Bacon
1958	\N	actor,producer,di...	You Should Have Left	You Should Have Left	
0	2020	\N	93	Horror,Mystery,Th...	actor \N
["Stetler"]					
	tt1270798	nm0000102	nm0000102	7.7	726015 Kevin Bacon
1958	\N	actor,producer,di...	X-Men: First Class	X: First Class	
0	2011	\N	131	Action,Sci-Fi	actor \N
["Sebastian Shaw"]					
	tt0822849	nm0000102	nm0000102	6.7	4277 Kevin Bacon
1958	\N	actor,producer,di...	Rails & Ties	Rails & Ties	
0	2007	\N	101	Drama	actor \N
["Tom Stark"]					
	tt0080761	nm0000102	nm0000102	6.4	158162 Kevin Bacon
1958	\N	actor,producer,di...	Friday the 13th	Friday the 13th	
0	1980	\N	95	Horror,Mystery,Th...	actor \N
["Jack"]					
	tt0094318	nm0000102	nm0000102	6.2	5946 Kevin Bacon
1958	\N	actor,producer,di...	White Water Summer	White Water Summer	
0	1987	\N	90	Adventure,Drama	actor \N
["Vic"]					
	tt0120303	nm0000102	nm0000102	6.2	2432 Kevin Bacon
1958	\N	actor,producer,di...	Telling Lies in A...	Telling Lies in A...	
0	1997	\N	101	Drama,Music	actor \N
["Billy Magic"]					
	tt0096094	nm0000102	nm0000102	5.9	13936 Kevin Bacon
1958	\N	actor,producer,di...	She's Having a Baby	She's Having a Baby	
0	1988	\N	106	Comedy,Drama,Romance	actor \N ["Jake
Briggs"]					
	tt0120890	nm0000102	nm0000102	6.6	131436 Kevin Bacon
1958	\N	actor,producer,di...	Wild Things	Wild Things	

```

0|      1998|      \N|      108| Crime,Drama,Mystery|      actor| \N|      ["Ray
Duquette"]|
+-----+-----+-----+-----+-----+-----+-----+
-----+-----+-----+-----+-----+-----+-----+
-----+-----+-----+-----+-----+-----+-----+
-----+
only showing top 20 rows

```

```

[ ]: graph = GraphFrame(vertices_df, edges_df)

# Kevin Bacon's ID
kevin_bacon_id = "nm0000102"

# Estimate shortest paths from Kevin Bacon to all other actors
shortest_paths = graph.shortestPaths(landmarks=[kevin_bacon_id])

# Extract and process the shortest path distances
shortest_paths = shortest_paths.select("id", col("distances").
    ↪getItem(kevin_bacon_id).alias("distance"))

# Handle unconnected actors by assigning a distance of 20
shortest_paths = shortest_paths.withColumn("distance", col("distance").
    ↪cast("int"))
shortest_paths = shortest_paths.na.fill(20, subset=["distance"])

# Cache the shortest paths DataFrame as it will be used later
shortest_paths.cache()

```

```
Out[34]: DataFrame[id: string, distance: int]
```

```

[ ]: # Q7 FINAL ANSWER
degree_summary = shortest_paths.groupBy("distance").count().orderBy("distance")

# Show the degree summary
degree_summary.show(15)

```

```

+-----+-----+
|distance|count|
+-----+-----+
|      0|    1|
|      1|  354|
|      2|14170|
|      3|58560|
|      4|42462|
|      5| 4842|
|      6|   510|
|      7|   56|

```

```
|      8|   20|
|      9|    3|
|     20|15330|
+-----+-----+
```

1.1.4 Exploring the data with RDD's

Using RDDs and (not dataframes) answer the following questions (if you loaded your data into spark in a dataframe you can convert to an RDD of rows easily using `.rdd`):

Q8 Movies can have multiple genres. Considering only titles of the type 'movie' what is the combination of genres that is the most popular (as measured by number of reviews). Hint: paired RDD's will be useful.

```
[ ]: from itertools import combinations
      from pyspark.sql.functions import col

      # Convert DataFrame to RDD and filter out movies with null numVotes
      movies_rdd = final_df.filter((col("titleType") == "movie") & (col("numVotes").
        ↪isNotNull())) .select("genres", "numVotes").rdd
```

```
[ ]: # Function to create genre combinations
      def create_genre_combinations(row):
          genres = row.genres.split(',')
          combinations = []
          for i in range(len(genres)):
              for j in range(i + 1, len(genres)):
                  combinations.append((f"{genres[i]},{genres[j]}", row.numVotes))
          return combinations
```

```
[ ]: # Create paired RDD with genre combinations and review counts
      genre_combinations_rdd = movies_rdd.flatMap(create_genre_combinations)
```

```
[ ]: # Sum the number of reviews for each genre combination
      genre_combination_counts = genre_combinations_rdd.reduceByKey(lambda a, b: a +
        ↪b)
```

```
[ ]: # Find the most popular genre combination
      most_popular_genre_combination = genre_combination_counts.max(lambda x: x[1])
```

```
[ ]: # Q8 FINAL ANSWER

      # Display the most popular genre combination and the number of reviews
      print(f"Most Popular Genre Combination: {most_popular_genre_combination[0]}.
        ↪Number of Reviews: {most_popular_genre_combination[1]}")
```

```
Most Popular Genre Combination: Action,Adventure. Number of Reviews: 4404499959
```

Q9 Movies can have multiple genres. Considering only titles of the type ‘movie’, and movies with more than 400 ratings, what is the combination of genres that has the highest **average movie rating** (you can average the movie rating for each movie in that genre combination). Hint: paired RDD’s will be useful.

```
[ ]: from itertools import combinations
      from pyspark.sql import Row

      # Convert DataFrames to RDDs
      title_basics_rdd = title_basics_df.rdd
      title_ratings_rdd = title_ratings_df.rdd

[ ]: # Filter for movies only and movies with more than 400 ratings
      movies_rdd = title_basics_rdd.filter(lambda row: row['titleType'] == 'movie')
      highly_rated_movies_rdd = title_ratings_rdd.filter(lambda row: row['numVotes'] >
      ↪ 400)

[ ]: # Join RDDs on tconst to get movie ratings
      movie_ratings_rdd = movies_rdd.map(lambda row: (row['tconst'], row)) \
      .join(highly_rated_movies_rdd.map(lambda row:
      ↪ (row['tconst'], (row['averageRating'], row['numVotes']))) \
      .map(lambda x: (x[1][0]['genres'], x[1][1]))

[ ]: # Function to create genre combinations and map to ratings
      def genre_combinations(row):
          genres = row[0]
          rating = row[1][0]
          if genres:
              genre_list = genres.split(',')
              for i in range(1, len(genre_list) + 1):
                  for combo in combinations(genre_list, i):
                      yield ('.'.join(sorted(combo)), (rating, 1))

[ ]: # Create genre combinations and map to ratings
      genre_combinations_rdd = movie_ratings_rdd.flatMap(genre_combinations)

[ ]: # Reduce by key to calculate the sum of ratings and the count for each genre
      ↪ combination
      genre_ratings_count_rdd = genre_combinations_rdd.reduceByKey(lambda a, b: (a[0]
      ↪ + b[0], a[1] + b[1]))

[ ]: # Calculate the average rating for each genre combination
      genre_avg_ratings_rdd = genre_ratings_count_rdd.mapValues(lambda v: v[0] / v[1])

[ ]: # Find the genre combination with the highest average rating
      highest_avg_rating_genre_combo = genre_avg_ratings_rdd.sortBy(lambda x: x[1],
      ↪ ascending=False).take(1)
```

```
[ ]: # Q9 FINAL ANSWER

# Show the highest average rating genre combination
for combo, avg_rating in highest_avg_rating_genre_combo:
    print(f"Highest Average Rating Genre Combination: {combo} with an average_
    ↳rating of {avg_rating}")
```

Highest Average Rating Genre Combination: Action,Documentary,Mystery with an average rating of 8.3

Q10 Movies can have multiple genres. What is the **individual genre** which is the most popular as measured by number of votes. Votes for multiple genres count towards each genre listed. Hint: flatmap and pairedRDD's will be useful here.

```
[ ]: # Convert DataFrames to RDDs
title_basics_rdd = title_basics_df.rdd
title_ratings_rdd = title_ratings_df.rdd
```

```
[ ]: # Filter for movies only
movies_rdd = title_basics_rdd.filter(lambda row: row['titleType'] == 'movie')
```

```
[ ]: # Join RDDs on tconst to get number of votes
movie_votes_rdd = movies_rdd.map(lambda row: (row['tconst'], row)) \
    .join(title_ratings_rdd.map(lambda row: (
    ↳(row['tconst'], row['numVotes']))) \
    .map(lambda x: (x[1][0]['genres'], x[1][1])))
```

```
[ ]: # Function to create individual genre records and map to votes
def explode_genres(row):
    genres = row[0]
    numVotes = row[1]
    if genres:
        genre_list = genres.split(',')
        for genre in genre_list:
            yield (genre, numVotes)
```

```
[ ]: # Create individual genre records and map to votes
genres_votes_rdd = movie_votes_rdd.flatMap(explode_genres)
```

```
[ ]: # Reduce by key to sum the number of votes for each genre
genre_votes_rdd = genres_votes_rdd.reduceByKey(lambda a, b: a + b)
```

```
[ ]: # Find the most popular genre by votes
most_popular_genre = genre_votes_rdd.sortBy(lambda x: x[1], ascending=False).
    ↳take(1)
```

```
[ ]: # Q10 FINAL ANSWER
```

```
# Show the most popular genre
for genre, votes in most_popular_genre:
    print(f"Most Popular Genre: {genre} with {votes} votes")
```

Most Popular Genre: Drama with 572360704 votes

1.2 Engineering the perfect cast

We have created a number of potential features for predicting the rating of a movie based on its cast. Use sparkML to build a simple linear model to predict the rating of a movie based on the following features:

1. The total number of movies in which the actors / actresses have acted (based on Q3)
2. The average pagerank of the cast in each movie (based on Q5)
3. The average outDegree of the cast in each movie (based on Q6)
4. The average value for for the cast of degrees of Kevin Bacon (based on Q7).

You will need to create a dataframe with the required features and label. Use a pipeline to create the vectors required by sparkML and apply the model. Remember to split your dataset, leave 30% of the data for testing, when splitting your data use the option seed=0.

Q11 Provide the coefficients of the regression and the accuracy of your model on that test dataset according to RSME.

```
[ ]: from pyspark.sql.functions import col, avg
      from pyspark.ml.feature import VectorAssembler
      from pyspark.ml.regression import LinearRegression
      from pyspark.ml import Pipeline
      from pyspark.ml.evaluation import RegressionEvaluator
```

```
[ ]: # Extract features
      # Q3: Total number of movies each actor has acted in
      actor_movie_counts_df = actor_movie_counts_rdd.toDF(["name_nconst",
      ↪ "total_movies"])
```

```
[ ]: # Q5: Average PageRank of the cast in each movie
      pagerank_df = pagerank_results.vertices.select(col("id").alias("name_nconst"),
      ↪ col("pagerank"))
```

```
[ ]: # Q6: Average outDegree of the cast in each movie
      # Extracting the first element of the struct which contains the numeric value
      outdegrees_df = outdegrees_with_names_rdd.map(lambda row: (row[0], row[1][0],
      ↪ row[1][1])).toDF(["name_nconst", "outdegree", "primaryName"])
```

```
[ ]: # Q7: Average degrees of separation from Kevin Bacon
      shortest_paths_df = shortest_paths.select(col("id").alias("name_nconst"),
      ↪ col("distance").alias("distance_to_kevin_bacon"))
```



```
[ ]: # Convert columns to numeric types where necessary
actor_movie_counts_df = actor_movie_counts_df.withColumn("total_movies",
    ↪col("total_movies").cast("double"))
pagerank_df = pagerank_df.withColumn("pagerank", col("pagerank").cast("double"))
outdegrees_df = outdegrees_df.withColumn("outdegree", col("outdegree").
    ↪cast("double"))
shortest_paths_df = shortest_paths_df.withColumn("distance_to_kevin_bacon",
    ↪col("distance_to_kevin_bacon").cast("double"))

[ ]: # Combine all features for each movie
features_df = filtered_df.select("tconst", "name_nconst").distinct()
features_df = features_df.join(actor_movie_counts_df, "name_nconst", "left")
features_df = features_df.join(pagerank_df, "name_nconst", "left")
features_df = features_df.join(outdegrees_df, "name_nconst", "left")
features_df = features_df.join(shortest_paths_df, "name_nconst", "left")

[ ]: # Aggregate features for each movie
aggregated_features_df = features_df.groupBy("tconst").agg(
    avg("total_movies").alias("avg_total_movies"),
    avg("pagerank").alias("avg_pagerank"),
    avg("outdegree").alias("avg_outdegree"),
    avg("distance_to_kevin_bacon").alias("avg_distance_to_kevin_bacon")
)

[ ]: # Fill null values with 0
aggregated_features_df = aggregated_features_df.na.fill(0)

[ ]: # Join with ratings to get the labels
ratings_df = final_df.select("tconst", "averageRating").distinct()
data = aggregated_features_df.join(ratings_df, "tconst")

[ ]: # Check for and handle any remaining null values
data = data.na.fill(0)

[ ]: # Prepare features and label
assembler = VectorAssembler(
    inputCols=["avg_total_movies", "avg_pagerank", "avg_outdegree",
    ↪"avg_distance_to_kevin_bacon"],
    outputCol="features"
)

[ ]: # Split data into training and test sets
train_data, test_data = data.randomSplit([0.7, 0.3], seed=0)

[ ]: # Define linear regression model
lr = LinearRegression(featuresCol="features", labelCol="averageRating")
```

```
[ ]: # Create pipeline
pipeline = Pipeline(stages=[assembler, lr])

[ ]: # Train the model
model = pipeline.fit(train_data)

[ ]: # Make predictions
predictions = model.transform(test_data)

[ ]: # Evaluate the model
evaluator = RegressionEvaluator(labelCol="averageRating",
    ↪ predictionCol="prediction", metricName="rmse")
rmse = evaluator.evaluate(predictions)

[ ]: # Get model coefficients
lr_model = model.stages[-1]
coefficients = lr_model.coefficients
intercept = lr_model.intercept

[ ]: # Q11 FINAL ANSWER

# Display the results
print(f"Coefficients: {coefficients}")
print(f"Intercept: {intercept}")
print(f"RMSE: {rmse}")
```

Coefficients: [-0.012659237327238589,-0.12336328797899417,0.012558636061121412,-0.0525044127818446]

Intercept: 4.168227402209818

RMSE: 2.9026865848009455

Q12 What score would your model predict for the 1997 movie Titanic.

```
[ ]: # Filter for the title "Titanic" and year 1997
titanic_df = title_basics_df.filter((col("primaryTitle") == "Titanic") &
    ↪ (col("startYear") == "1997"))

# Show the filtered result
titanic_df.show()

# Collect the ID for Titanic movie
titanic_id = titanic_df.select("tconst").collect()[0][0]
```

```
+-----+-----+-----+-----+-----+-----+-----+-----+
+-----+
|  tconst|titleType|primaryTitle|originalTitle|isAdult|startYear|endYear|runtimeMinutes|genres|
+-----+-----+-----+-----+-----+-----+-----+-----+
+-----+
+-----+-----+-----+-----+-----+-----+-----+-----+
```

tt0120338	movie	Titanic	Titanic	0	1997	\N
194	Drama,Romance					
tt0594950	tvEpisode	Titanic	Titanic	0	1997	\N
\N	Documentary,Short					
tt5722820	tvEpisode	Titanic	Titanic	0	1997	\N
\N	Documentary,News					

-----+-----+-----+-----+-----+-----+-----+-----+-----
-----+-----+-----+-----+-----+-----+-----+-----+-----

```
[ ]: # Filter the data for Titanic
titanic_features = data.filter(col("tconst") == titanic_id)

[ ]: # Step 3: Use the trained model to predict Titanic's rating
titanic_prediction = model.transform(titanic_features)

[ ]: # Q12 FINAL ANSWER

# Extract and print the predicted rating
predicted_rating = titanic_prediction.select("prediction").collect()[0][0]
print(f"Predicted rating for Titanic (1997): {round(predicted_rating,2)}")
```

Predicted rating for Titanic (1997): 5.88

Q13 Create dummy variables for each of the top 10 movie genres for Q10. These variable should have a value of 1 if the movie was rated with that genre and 0 otherwise. For example the 1997 movie Titanic should have a 1 in the dummy variable column for Romance, and a 1 in the dummy variable column for Drama, and 0's in all the other dummy variable columns.

Does adding these variable to the regression improve your results? What is the new RMSE and predicted rating for the 1997 movie Titanic.

```
[ ]: from pyspark.sql.functions import col, when, split, sum as _sum
from pyspark.ml.feature import VectorAssembler
from pyspark.ml.regression import LinearRegression
from pyspark.ml import Pipeline
from pyspark.ml.evaluation import RegressionEvaluator

[ ]: # Filter for movies only
movies_df = final_df.filter(col("titleType") == "movie")

[ ]: # Explode genres
movies_with_genres_df = movies_df.withColumn("genre",
↳explode(split(col("genres"), ",")))

[ ]: # Calculate total votes for each genre
genre_votes_df = movies_with_genres_df.groupBy("genre").agg(_sum("numVotes").
↳alias("total_votes"))
```

```
[ ]: # Sort by total votes and select top 10 genres
top_10_genres_df = genre_votes_df.orderBy(col("total_votes").desc()).limit(10)
top_10_genres = [row['genre'] for row in top_10_genres_df.collect()]
```

```
[ ]: # Display the top 10 genres
print("Top 10 genres by number of votes:")
for genre in top_10_genres:
    print(genre)
```

Top 10 genres by number of votes:

- Drama
- Action
- Comedy
- Adventure
- Crime
- Thriller
- Sci-Fi
- Romance
- Mystery
- Horror

```
[ ]: # Create dummy variables for the top 10 genres in final_df
for genre in top_10_genres:
    final_df = final_df.withColumn(f"genre_{genre}", when(col("genres").
        ↳contains(genre), 1).otherwise(0))
```

```
[ ]: # Q3: Total number of movies each actor has acted in
actor_movie_counts_df = actor_movie_counts_rdd.toDF(["name_nconst",
    ↳"total_movies"])
```

```
[ ]: # Q5: Average PageRank of the cast in each movie
pagerank_df = pagerank_results.vertices.select(col("id").alias("name_nconst"),
    ↳col("pagerank"))
```

```
[ ]: # Q6: Average outDegree of the cast in each movie
outdegrees_df = outdegrees_with_names_rdd.map(lambda row:
    ↳Row(name_nconst=row[0], outdegree=row[1][0])).toDF()
```

```
[ ]: # Q7: Average degrees of separation from Kevin Bacon
shortest_paths_df = shortest_paths.select(col("id").alias("name_nconst"),
    ↳col("distance").alias("distance_to_kevin_bacon"))
```

```
[ ]: # Combine all features for each movie
features_df = filtered_df.select("tconst", "name_nconst").distinct()
features_df = features_df.join(actor_movie_counts_df, "name_nconst", "left")
features_df = features_df.join(pagerank_df, "name_nconst", "left")
features_df = features_df.join(outdegrees_df, "name_nconst", "left")
features_df = features_df.join(shortest_paths_df, "name_nconst", "left")
```

```
[ ]: # Aggregate features for each movie
aggregated_features_df = features_df.groupBy("tconst").agg(
    avg("total_movies").alias("avg_total_movies"),
    avg("pagerank").alias("avg_pagerank"),
    avg("outdegree").alias("avg_outdegree"),
    avg("distance_to_kevin_bacon").alias("avg_distance_to_kevin_bacon")
)

[ ]: # Fill null values with 0
aggregated_features_df = aggregated_features_df.na.fill(0)

[ ]: # Join with ratings to get the labels
ratings_df = final_df.select("tconst", "averageRating").distinct()
data = aggregated_features_df.join(ratings_df, "tconst")

[ ]: # Join with dummy variables for genres
for genre in top_10_genres:
    genre_col = f"genre_{genre}"
    genre_df = final_df.select("tconst", genre_col).distinct()
    data = data.join(genre_df, "tconst", "left").na.fill(0)

[ ]: # Check for and handle any remaining null values
data = data.na.fill(0)

[ ]: from pyspark.ml.feature import VectorAssembler, StandardScaler

# Prepare feature set
assembler = VectorAssembler(
    inputCols=["avg_total_movies", "avg_pagerank", "avg_outdegree",
    ↪ "avg_distance_to_kevin_bacon"] + [f"genre_{genre}" for genre in
    ↪ top_10_genres],
    outputCol="features"
)
scaler = StandardScaler(inputCol="features", outputCol="scaledFeatures",
    ↪ withStd=True, withMean=False)

[ ]: # Split data into training and test sets
train_data, test_data = data.randomSplit([0.7, 0.3], seed=0)

[ ]: # Define linear regression model
lr = LinearRegression(featuresCol="scaledFeatures", labelCol="averageRating")

[ ]: # Create pipeline
pipeline = Pipeline(stages=[assembler, scaler, lr])

[ ]: model = pipeline.fit(train_data)
```

```
[ ]: # Make predictions
predictions = model.transform(test_data)
```

```
[ ]: # Evaluate the model
evaluator = RegressionEvaluator(labelCol="averageRating",
    ↪ predictionCol="prediction", metricName="rmse")
rmse = evaluator.evaluate(predictions)
```

```
[ ]: # Q13 PART 1 FINAL ANSWER

# Display the RMSE
print(f"New RMSE with genre dummy variables: {round(rmse,4)}")
print(f"This is an improvement in relation to the model without genre dummies")
```

New RMSE with genre dummy variables: 2.792

This is an improvement in relation to the model without genre dummies

```
[ ]: # Predict the rating for Titanic using the updated model
titanic_id = final_df.filter(col("primaryTitle") == "Titanic").select("tconst").
    ↪ distinct().first()[0]
```

```
[ ]: # Extract features for Titanic
titanic_features_df = data.filter(col("tconst") == titanic_id)
```

```
[ ]: titanic_features_df.drop("averageRating")
```

```
Out[110]: DataFrame[tconst: string, avg_total_movies: double, avg_pagerank:
double, avg_outdegree: double, avg_distance_to_kevin_bacon: double, genre_Drama:
int, genre_Action: int, genre_Comedy: int, genre_Adventure: int, genre_Crime:
int, genre_Thriller: int, genre_Sci-Fi: int, genre_Romance: int, genre_Mystery:
int, genre_Horror: int]
```

```
[ ]: # Fill null values with 0
titanic_features_df = titanic_features_df.na.fill(0)
```

```
[ ]: # Prepare features for prediction
titanic_features_vector = assembler.transform(titanic_features_df)
```

```
[ ]: # Ensure no conflicting columns for prediction
titanic_features_vector = titanic_features_vector.drop("features")
```

```
[ ]: # Assemble features for Titanic again
titanic_assembler = VectorAssembler(
    inputCols=["avg_total_movies", "avg_pagerank", "avg_outdegree",
    ↪ "avg_distance_to_kevin_bacon"] + [f"genre_{genre}" for genre in
    ↪ top_10_genres],
    outputCol="features"
)
```

```
[ ]: titanic_features_vector = titanic_assembler.transform(titanic_features_df)

[ ]: # Making sure there is no conflict as we were getting an error in the next cell
    ↳ saying "column features already exists"
titanic_features_vector = titanic_features_vector.drop("features")

[ ]: # Make prediction for Titanic
titanic_prediction = model.transform(titanic_features_vector)

[ ]: # Extract and print the predicted rating
predicted_rating = titanic_prediction.select("prediction").collect()

[ ]: # Q13 PART 2 FINAL ANSWER
print(f"Predicted rating for Titanic (1997): {predicted_rating[0][0]}")
```

Predicted rating for Titanic (1997): 5.187839660493992

Q14 - Open Question: Improve your model by testing different machine learning algorithms, using hyperparameter tuning on these algorithms, changing the included features. What is the RMSE of your final model and what rating does it predict for the 1997 movie Titanic.

```
[ ]: from pyspark.sql.functions import col, avg, when, lit
from pyspark.ml.feature import VectorAssembler, StandardScaler
from pyspark.ml.regression import LinearRegression, RandomForestRegressor,
    ↳ GBTRRegressor
from pyspark.ml import Pipeline
from pyspark.ml.evaluation import RegressionEvaluator
from pyspark.ml.tuning import CrossValidator, ParamGridBuilder

[ ]: # Convert RDDs to DataFrames
actor_movie_counts_df = actor_movie_counts_rdd.toDF(["name_nconst",
    ↳ "total_movies"])
pagerank_df = pagerank_results.vertices.select(col("id").alias("name_nconst"),
    ↳ col("pagerank"))
outdegrees_df = outdegrees_with_names_rdd.map(lambda row:
    ↳ Row(name_nconst=row[0], outdegree=row[1][0])).toDF()
shortest_paths_df = shortest_paths.select(col("id").alias("name_nconst"),
    ↳ col("distance").alias("distance_to_kevin_bacon"))

[ ]: # Combine all features for each movie
features_df = filtered_df.select("tconst", "name_nconst").distinct()
features_df = features_df.join(actor_movie_counts_df, "name_nconst", "left")
features_df = features_df.join(pagerank_df, "name_nconst", "left")
features_df = features_df.join(outdegrees_df, "name_nconst", "left")
features_df = features_df.join(shortest_paths_df, "name_nconst", "left")

[ ]: # Combine all features for each movie
features_df = filtered_df.select("tconst", "name_nconst").distinct()
```

```

features_df = features_df.join(actor_movie_counts_df, "name_nconst", "left")
features_df = features_df.join(pagerank_df, "name_nconst", "left")
features_df = features_df.join(outdegrees_df, "name_nconst", "left")
features_df = features_df.join(shortest_paths_df, "name_nconst", "left")

```

```

[ ]: # Aggregate features for each movie
aggregated_features_df = features_df.groupBy("tconst").agg(
    avg("total_movies").alias("avg_total_movies"),
    avg("pagerank").alias("avg_pagerank"),
    avg("outdegree").alias("avg_outdegree"),
    avg("distance_to_kevin_bacon").alias("avg_distance_to_kevin_bacon")
)

```

```

[ ]: # Fill null values with 0
aggregated_features_df = aggregated_features_df.na.fill(0)

```

```

[ ]: # Join with ratings to get the labels
ratings_df = final_df.select("tconst", "averageRating").distinct()
data = aggregated_features_df.join(ratings_df, "tconst")

```

```

[ ]: # Check for and handle any remaining null values
data = data.na.fill(0)

```

```

[ ]: # Prepare feature set
assembler = VectorAssembler(
    inputCols=["avg_total_movies", "avg_pagerank", "avg_outdegree", ↵
    ↵"avg_distance_to_kevin_bacon"],
    outputCol="features4"
)
scaler_model = StandardScaler(inputCol="features4", ↵
    ↵outputCol="scaledFeatures4", withStd=True, withMean=False)

```

```

[ ]: # Split data into training and test sets
train_data, test_data = data.randomSplit([0.7, 0.3], seed=0)

```

```

[ ]: # Define models
lr = LinearRegression(featuresCol="scaledFeatures4", labelCol="averageRating")
rf = RandomForestRegressor(featuresCol="scaledFeatures4", ↵
    ↵labelCol="averageRating")
gbt = GBTRegressor(featuresCol="scaledFeatures4", labelCol="averageRating")

```

```

[ ]: # Define parameter grids for hyperparameter tuning
paramGridLR = ParamGridBuilder().addGrid(lr.regParam, [0.01, 0.1, 0.5]).build()
paramGridRF = ParamGridBuilder().addGrid(rf.numTrees, [20, 50, 100]).build()
paramGridGBT = ParamGridBuilder().addGrid(gbt.maxIter, [10, 20, 50]).build()

```



```
[ ]: # Define evaluators
evaluator = RegressionEvaluator(labelCol="averageRating",
    predictionCol="prediction", metricName="rmse")

[ ]: # Cross-validation for each model
crossvalLR = CrossValidator(estimator=Pipeline(stages=[assembler, scaler_model,
    lr]),
    estimatorParamMaps=paramGridLR,
    evaluator=evaluator,
    numFolds=5)
crossvalRF = CrossValidator(estimator=Pipeline(stages=[assembler, scaler_model,
    rf]),
    estimatorParamMaps=paramGridRF,
    evaluator=evaluator,
    numFolds=5)
crossvalGBT = CrossValidator(estimator=Pipeline(stages=[assembler,
    scaler_model, gbt]),
    estimatorParamMaps=paramGridGBT,
    evaluator=evaluator,
    numFolds=5)

[ ]: # Fit models
cvModelLR = crossvalLR.fit(train_data)
cvModelRF = crossvalRF.fit(train_data)
cvModelGBT = crossvalGBT.fit(train_data)

[ ]: # Evaluate models
predictionsLR = cvModelLR.transform(test_data)
predictionsRF = cvModelRF.transform(test_data)
predictionsGBT = cvModelGBT.transform(test_data)

rmseLR = evaluator.evaluate(predictionsLR)
rmseRF = evaluator.evaluate(predictionsRF)
rmseGBT = evaluator.evaluate(predictionsGBT)

[ ]: print(f"RMSE for Linear Regression: {rmseLR}")
print(f"RMSE for Random Forest: {rmseRF}")
print(f"RMSE for Gradient Boosted Trees: {rmseGBT}")

RMSE for Linear Regression: 2.902608241094834
RMSE for Random Forest: 2.8266434489955774
RMSE for Gradient Boosted Trees: 2.7850901282218037

[ ]: # Q14 PART 1 FINAL ANSWER

# Choose the best model
print(f"Our best model is Gradient Boosted Trees with an RMSE of {rmseGBT}.")
```

Our best model is Gradient Boosted Trees with an RMSE of 2.7850901282218037.

```
[ ]: # Predict the rating for Titanic using the best model

# Extract features for Titanic
titanic_features_df = filtered_df.filter(col("tconst") == titanic_id).
    ↪select("tconst", "name_nconst").distinct()

[ ]: # Perform the joins
titanic_features_df = titanic_features_df.join(actor_movie_counts_df,
    ↪"name_nconst", "left")
titanic_features_df = titanic_features_df.join(pagerank_df, "name_nconst",
    ↪"left")
titanic_features_df = titanic_features_df.join(outdegrees_df, "name_nconst",
    ↪"left")
titanic_features_df = titanic_features_df.join(shortest_paths_df,
    ↪"name_nconst", "left")

# Aggregate features for Titanic
titanic_aggregated_features_df = titanic_features_df.groupBy("tconst").agg(
    avg("total_movies").alias("avg_total_movies"),
    avg("pagerank").alias("avg_pagerank"),
    avg("outdegree").alias("avg_outdegree"),
    avg("distance_to_kevin_bacon").alias("avg_distance_to_kevin_bacon")
)

[ ]: # Add dummy variables for Titanic
for genre in top_10_genres:
    titanic_aggregated_features_df = titanic_aggregated_features_df.
        ↪withColumn(f"genre_{genre}", lit(1) if genre in ["Romance", "Drama"] else
        ↪lit(0))

# Fill null values with 0
titanic_aggregated_features_df = titanic_aggregated_features_df.na.fill(0)

[ ]: # Prepare feature set
assembler = VectorAssembler(
    inputCols=["avg_total_movies", "avg_pagerank", "avg_outdegree",
        ↪"avg_distance_to_kevin_bacon"] + [f"genre_{genre}" for genre in
        ↪top_10_genres],
    outputCol="raw_features"
)

# Assemble the features
assembled_data = assembler.transform(titanic_aggregated_features_df)
```

```
[ ]: # Prepare features for prediction
# Scale the features
scaler3 = StandardScaler(inputCol="raw_features", outputCol="scaledFeatures4",
    ↳withStd=True, withMean=False)

[ ]: # Making sure there is no conflict as we were getting an error in the next cell
    ↳saying "column features4 already exists"
titanic_features_vector = titanic_features_vector.drop("scaledFeatures4")

[ ]: # Make prediction for Titanic
titanic_prediction = cvModelGBT.transform(titanic_features_vector)

[ ]: # Extract and print the predicted rating
predicted_rating = titanic_prediction.select("prediction").collect()

[ ]: # Q14 FINAL ANSWER PART 2

# Display the prediction for Titanic
print(f"Predicted Rating for Titanic with the best model (Gradient Boosted
    ↳Trees): {round(predicted_rating[0][0],2)}")
```

Predicted Rating for Titanic with the best model (Gradient Boosted Trees): 4.77

2 Statement of Authorship

André Filipe Silva - 20230972 João Gonçalves - 20230560